

# AI Legal Consultant

Arshdeep Singh, Soumya Bhate, Yatharth Kumar





# Whoa!!

This AI assistant provides general legal information only.  
It does NOT constitute legal advice. For specific legal matters, please consult a licensed attorney.

# Table of Contents

**01** Project Overview

**02** Knowledge Base and Chunking

**03** Embedding and Vector Search

**04** RAG Answer Generation

**05** User Interface and Experience

**06** Future Enhancements



01

# Project Overview

An abstract of the legal AI chatbot





# Project Overview

Purpose: Build a chatbot that answers criminal law questions using a RAG pipeline

Key Features-

- Domain-specific knowledge base (Eg- Criminal, Contract, Real estate Law etc)
- MongoDB vector search with embeddings
- Open AI powered LLM response generation
- Text-to-speech (TTS) and Speech to text (STT) via Eleven Labs
- Upload your legal documents and ask away!
- Downloadable consultations and audio device settings





02

# Knowledge Base & Chunking

Using MongoDB to create a KB for a RAG style bot

# KB and Chunking

- PDF Upload & Extraction using PyPDF2 to extract text from Law Books and legal documents.
- Chunking sentences of this extracted text with overlap (size=500, overlap=100) to provide context across chunks
- Metadata Stored in MongoDB with chunk\_id, pdf\_name, chunk\_size, created\_at
- KB in MongoDB decreases the risk of hallucinations faced by most transformer models



# Why do we need a KB?

- We know that laws are getting updated almost every day and it is not possible for LLMs to keep up to it.
  - Building a rag style chatbot makes it feasible for LLM to get richer context and help you make better decisions!
  - The screenshot shows chunks of a criminal law pdf. The LLM can now answer your questions from your uploaded text with this Criminal Law Pdf as context.

03

# Embedding & Vector Search

Discussing about how we are embedding the chunk  
of texts and using vector search



# Embedding & Vector Search

## Embedding Model

Via Sentence Transformers

## Vector Search

Retrieves top-k relevant chunks based on query embedding

all-miniLM-L6  
-v2

MongoDB

\$vectorSearch

Relevant Information from  
the KB

## Storage

Embeddings stored in a NoSQL DB using bulk write

## Output

User query will be the input and vector search will make sure it provides relevant information from the selected KB (streamlit Legal Area option)

04

# RAG Answer Generation

Explains how KB is being used to answer questions



# RAG Answer Generation

## Search Function:

- Retrieves top 3 (customizable) relevant chunks using semantic similarity

## Prompt Construction:

- Injects retrieved context into Groq prompt
- Uses openai/gpt-oss-120b for response generation

## Guardrails:


- Bot only responds to legal queries
- Prompt role curated for “legal advisor”
- Disclaimer embedded in prompt and UI



# 05

# User Interface and Experience

Talks about the Streamlit UI with python and the interaction flow



# UI and Interaction Flow

## Streamlit UI Features:

- Select your Legal Area that you have questions about
- Criminal Law KB toggle (MongoDB-backed, will upload more documents)
- Audio device settings and TTS playback (TTS not added yet)
- Downloadable consultation as .txt
- Conversation history maintained

## Interaction Flow:

- “Ask Your Legal Question” button / Upload your document
- Input field with placeholder guidance
- Legal disclaimer prominently displayed



06

# Future Enhancements

Milestones to Complete



# Features to add



Multi-domain KB support (Civil, Corporate, etc.)



Upload PDF/ document option for the users with personal details redacted



Real-time audio playback (TTS)



Model fine-tuning for legal tone and accuracy and improve UI



# Thanks!

**Does anyone have any questions?**

